

Data Parallel Execution Model

John H. Osorio Ríos



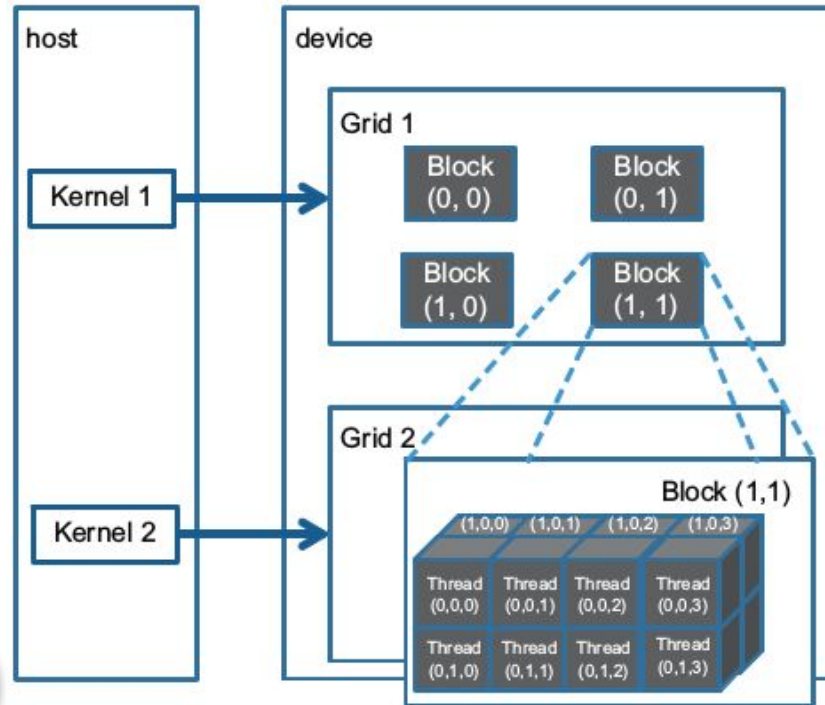
Cuda Thread Organization (1/2)

```
dim3 dimGrid(ceil(n/256.0), 1, 1);  
dim3 dimBlock(256, 1, 1);  
vecAddKernel << <dimGrid, dimBlock>> > (...);
```

```
dim3 dimBlock(2, 2, 1);  
dim3 dimGrid(4, 2, 2);  
KernelFunction << <dimGrid, dimBlock>> > (...);
```




Cuda Thread Organization (2/2)




- **blockDim**
- **threadIdx**
- **blockIdx**
- **gridDim**



Mapping Threads to Multidimensional Data (1/4)




16x16 blocks




- 76 x 62 Pixels
- 5 x blocks
- 4 y blocks
- 16 x 16 threads in a block



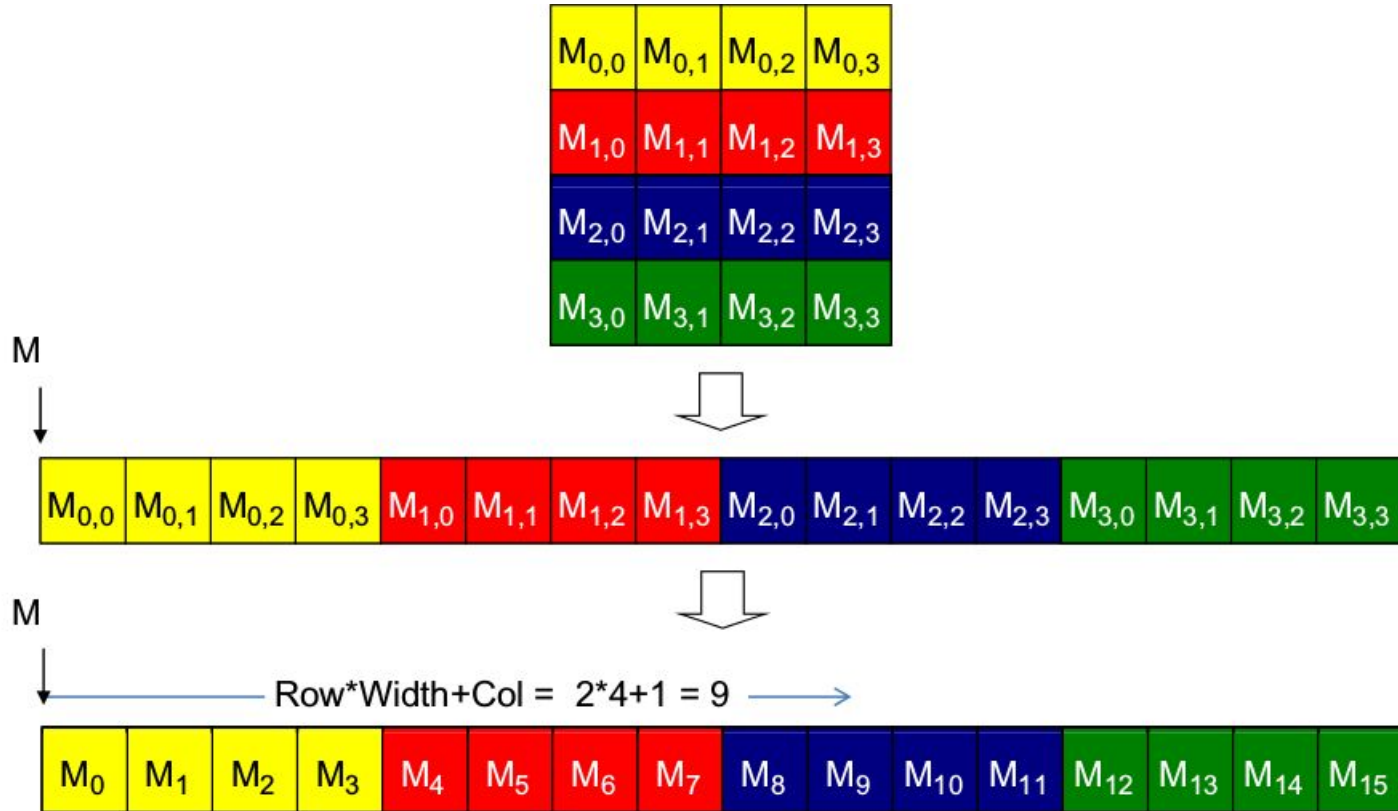
Mapping Threads to Multidimensional Data (2/4)



16x16 block



Mapping Threads to Multidimensional Data (3/4)



Mapping Threads to Multidimensional Data (4/4)

```
__global__ void PictureKernell(float* d_Pin, float* d_Pout, int n, int m) {  
  
    // Calculate the row # of the d_Pin and d_Pout element to process  
    int Row = blockIdx.y*blockDim.y + threadIdx.y;  
  
    // Calculate the column # of the d_Pin and d_Pout element to process  
    int Col = blockIdx.x*blockDim.x + threadIdx.x;  
  
    // each thread computes one element of d_Pout if in range  
    if ((Row < m) && (Col < n)) {  
        d_Pout[Row*n+Col] = 2*d_Pin[Row*n+Col];  
    }  
  
}
```

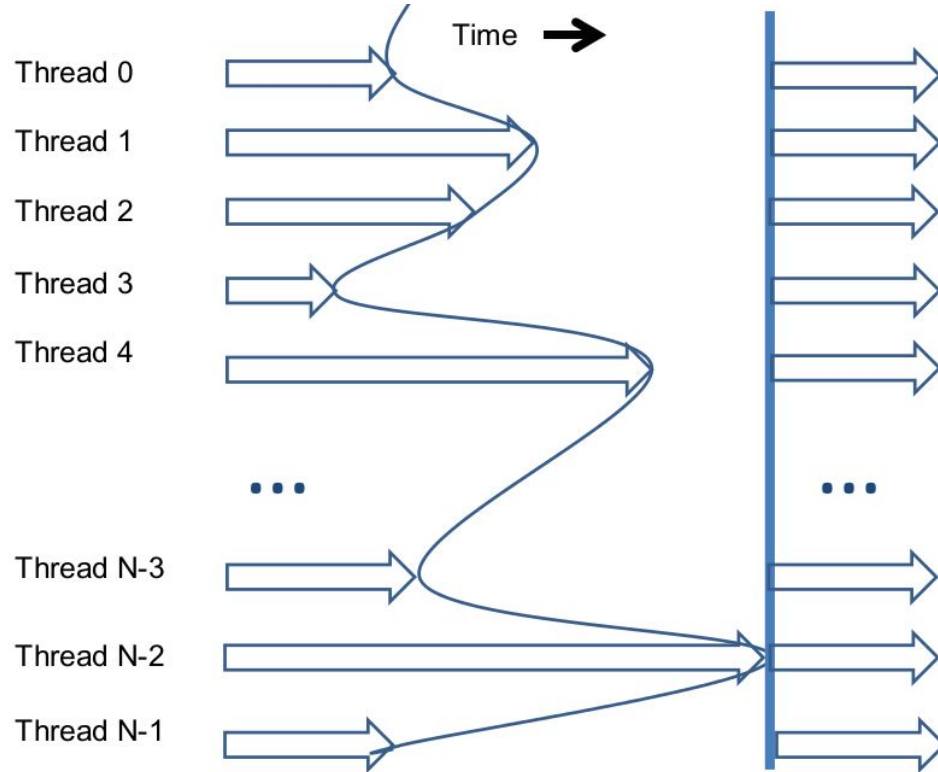


Matrix-Matrix Multiplication

```
__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width) {  
  
    // Calculate the row index of the d_Pelement and d_M  
    int Row = blockIdx.y*blockDim.y+threadIdx.y;  
  
    // Calculate the column index of d_P and d_N  
    int Col = blockIdx.x*blockDim.x+threadIdx.x;  
  
    if ((Row < Width) && (Col < Width)) {  
        float Pvalue = 0;  
        // each thread computes one element of the block sub-matrix  
        for (intk = 0; k < Width; ++k) {  
            Pvalue += d_M[Row*Width+k]*d_N[k*Width+Col];  
        }  
        d_P[Row*Width+Col] = Pvalue;  
    }  
}
```



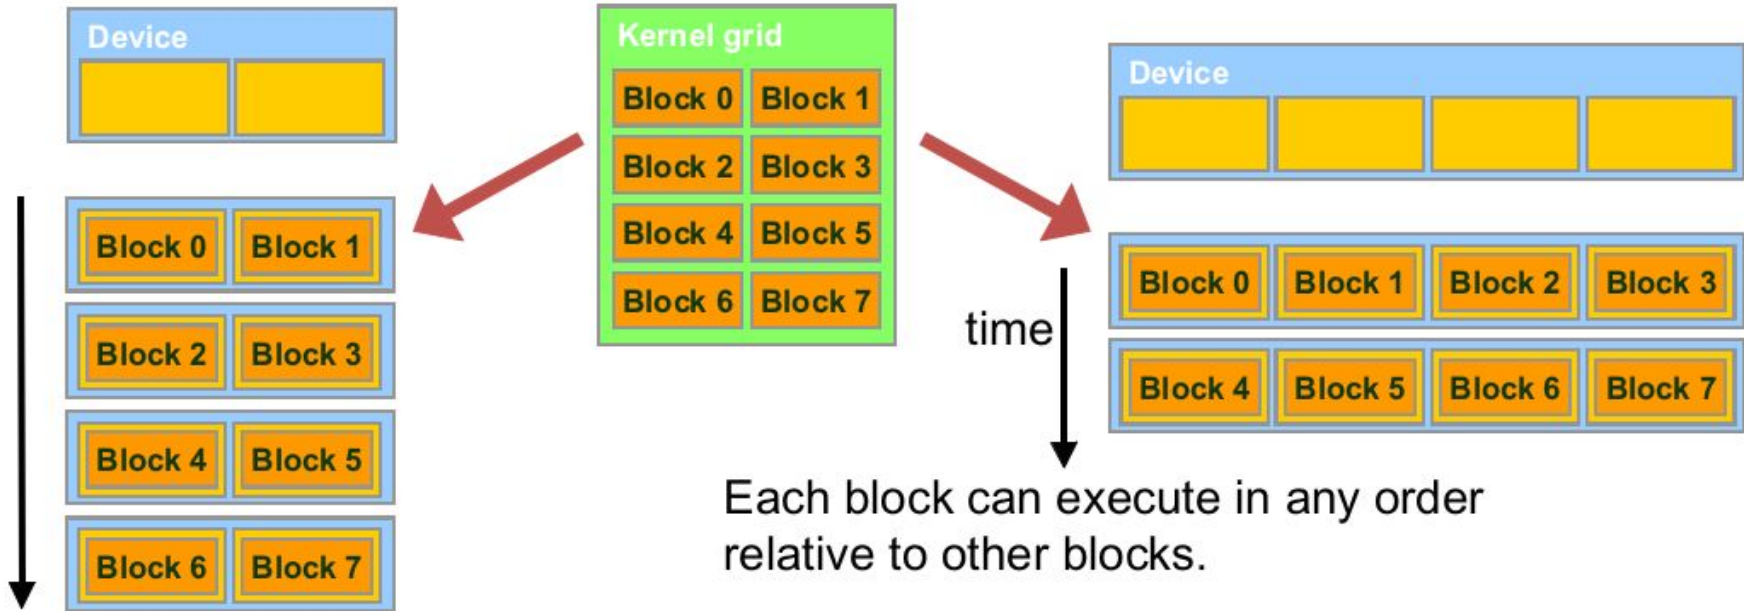
Synchronization and Transparent Scalability (1/2)




- `__syncthreads()`



Synchronization and Transparent Scalability (2/2)



Assigning Resources to Blocks



- 8 block to each SM
- 1536 threads to each SM
- 6 blocks of 256 threads
- 3 blocks of 512
- **12 blocks of 128 threads**

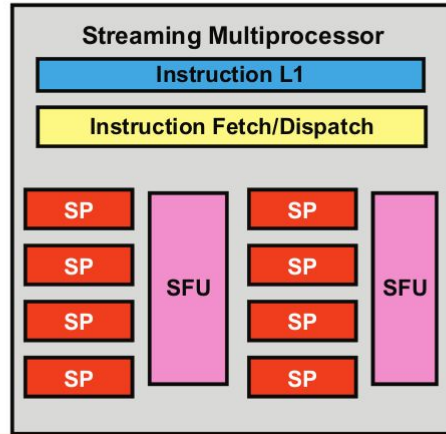



Querying Device Properties

```
int dev_count;  
cudaGetDeviceCount( &dev_count);  
  
cudaDeviceProp dev_prop;  
for (I = 0; i < dev_count; i++) {  
    cudaGetDeviceProperties( &dev_prop, i);  
    // decide if device has sufficient resources and capabilities  
}
```



Thread Scheduling and Latency Tolerance



- **Assume a CUDA device:**
 - 8 blocks per SM
 - 1024 threads per SM
 - 512 threads per block
- **For Matrix Multiplication:**
 - 8x8 thread blocks ?
 - 16x16 thread blocks ?
 - 32x32 thread blocks ?



THANKS

john@sirius.utp.edu.co

